

Attribution gating — scaling components by their attribution to the token

Attribution is the GIM sensor the decomposition was built from: a single modified backward ($\tau=2$ softmax Jacobian, scaling/4 on Q and K, halved gate/up gradients, frozen RMSNorm statistics) against a PRE-softmax objective that puts one unit of reward on each ground-truth next-token logit — not plain gradients and not integrated gradients. A component's score at token t is $\sum |\partial \text{Reward} / \partial W \cdot W|$ over the entries it owns. Weights are shared across positions, so this is one weight rebuild and one forward per position: 24 positions $\times$ 46 configurations in 168 s.

