


How components interact at runtime

Co-activation is the frozen fingerprint posterior, the same quantity the streaming assignment uses. Lift = $\text{Coact}(a,b) / p(a)p(b)$ — how much more two components fire together than their marginals predict. Synergy is measured by ablating whole components: $\text{dCE}(\text{both}) - \text{dCE}(a) - \text{dCE}(b)$, relative to the main effects. Both arms are drawn from the SAME pool of components (those whose single ablation is measurable), so main effects are matched and only the pairing differs. Components that fire together need each other; components that never fire together partly substitute for each other.

The runtime graph is diffuse, not modular



Co-firing predicts the sign of causal interaction

