



Editing induction: two components, two different failures

Each component's owned weight mass is scaled by α in all 112 matrices. The induction score is how much attention flows from a repeated token to the token that followed it the first time — the mechanism itself, measured independently of the loss.



Do the induction heads still look in the right place?



Mechanism vs readout

