Which attribution sensor yields the most ablatable decomposition? (67M Pile target, C=256)

Everything downstream of the sensor is held fixed — PCA → spherical k-means on token rows, softpart bank, minimality metric — only the sensor varies, applied end-to-end as production applies GIM. IG integrates along the WEIGHT-SCALING path this codebase already uses (--ig_k): each decomposable matrix scaled by α=k/K, a zero-weight baseline. Five seeds, each sharing pilot positions and k-means init across all nine sensors, so the tests are paired. Plain grad×act beats the production sensor by +11.8 components (t=4.00, p=0.016); integrating makes it monotonically worse (EAP−IG5 = +12.4, p=0.004); GIM's τ-softmax fix alone already matches EAP.