


Editing one component family to unlearn hazardous biology

WMDP-bio scored on 823 questions the edit never saw; the edit is fit on 300 others. Only the components' own weight mass is scaled — 0.0024% of the model at k=4 — with one scalar per matrix per component. Unlike the German edit there is no free corner: past roughly a third of the removable knowledge, general capability starts coming out with it. Adding components buys a strictly better trade-off, not a clean one.

Every configuration swept



At the operating point that keeps MMLU (≤ 0.5 pt)

